

permute: A Python Package for Randomization Inference

Kellie Ottoboni

Department of Statistics, UC Berkeley
Berkeley Institute for Data Science

June 16, 2016
ISNPS, Avignon



University of California, Berkeley
DEPARTMENT OF STATISTICS



Collaborators



Jarrod Millman
jarrodmillman



Philip B. Stark
pbstark



**Stefan van der
Walt**
stefanv

Outline

1 Introduction

2 Examples

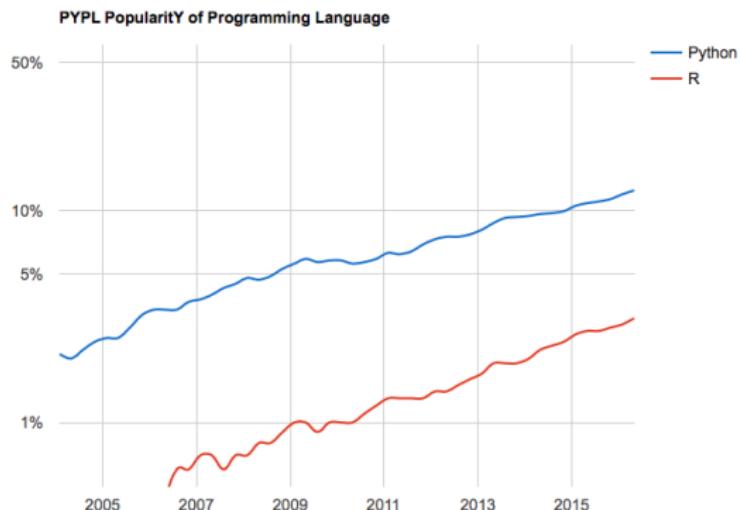
- Gender bias in teaching evaluations
- Inter-rater reliability

3 The role of software development in Statistics

Python

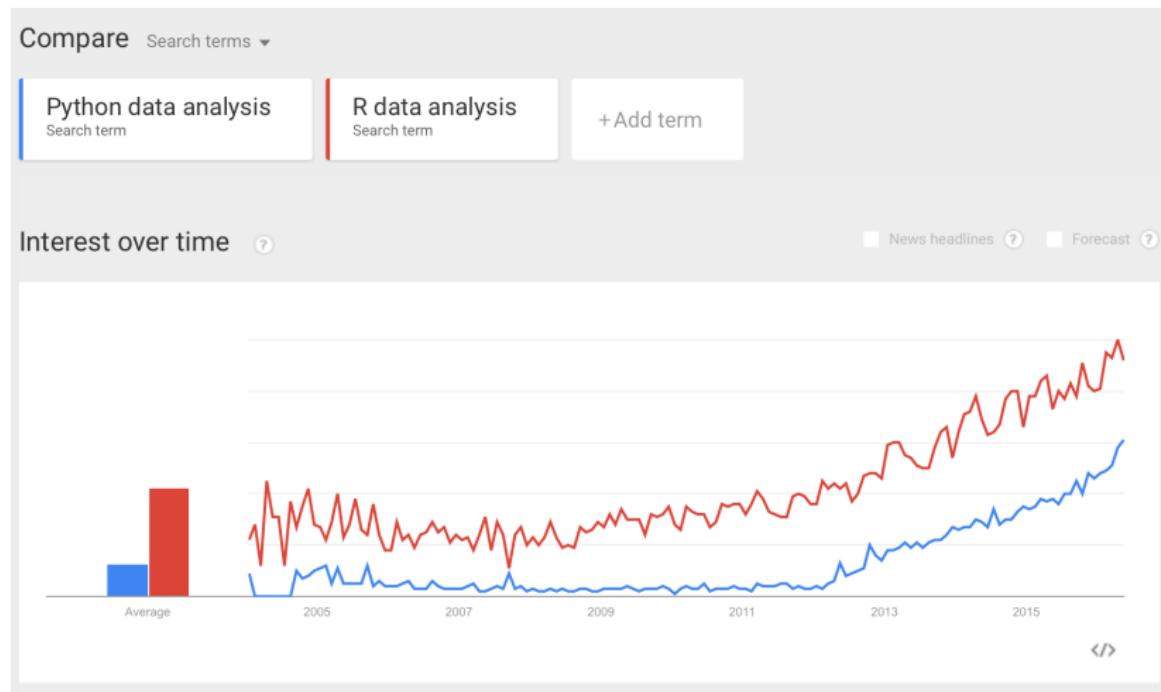
Python is gaining popularity for doing data analysis.

- General purpose language with “batteries included”
- Popular for a variety of scientific applications



Python

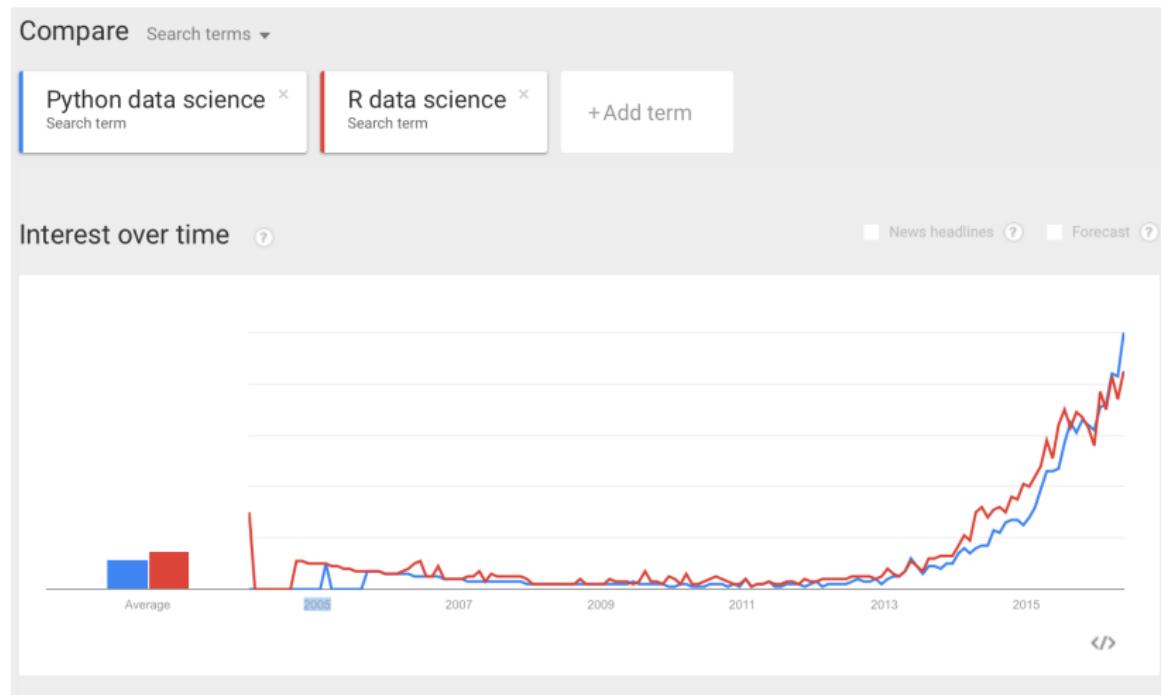
Google trends on May 22, 2016



Keyword: data analysis

Python

Google trends on May 22, 2016



Keyword: data science

Python for teaching Statistics



Data Science 8, Spring 2016 at UC Berkeley

Permutation tests

- Fisher [1935] introduced permutation tests for randomized experiments
- TO DO: ADD MORE REFS!
- Rely on assumptions about randomization or exchangeability, rather than parametric assumptions, IID sampling, etc.

James Bradley [1968]

“[a] corresponding parametric test is valid only to the extent that it results in the same statistical decision [as the randomization test].”

Permutation tests

R has several packages for randomization inference.

- `ri`
- `RIttools`
- `coin`
- `perm`

In Python, statistics packages are limited.

- `pandas`
- `scipy.stats`
- `StatsModels`
- `scikit-learn`

Download permute!

Permutation tests and confidence sets

[build](#) passing [coverage](#) 99%

Permutation tests and confidence sets for a variety of nonparametric testing and estimation problems, for a variety of randomization designs.

- **Website (including documentation):** <http://statlab.github.io/permute>
- **Mailing list:** <http://groups.google.com/group/permute>
- **Source:** <https://github.com/statlab/permute>
- **Bug reports:** <https://github.com/statlab/permute/issues>

Installation from binaries

```
$ pip install permute
```

Teaching Evaluations

Student evaluations of teachers (SET) are used to

- Quantify teaching effectiveness
- Compare instructors across courses
- Make hiring, firing, and promotion decisions

Are SET a valid measure of teaching effectiveness?

Teaching evaluations

In Boring et al. [2016], we reanalyzed data from MacNell et al. [2014].

- Students were randomized to 4 online sections of a course.
- In two sections, the instructors swapped identities.
- Was the instructor who identified as female rated lower on average?

Neyman-Rubin model, generalized

Student i is represented by a ticket with 4 numbers, their response to each “treatment.”

$$r_{ijk} = \begin{aligned} &\text{SET given by student } i \text{ to instructor } j \\ &\text{when they appear to have gender } k \\ i = 1, \dots, N; \quad j = 1, 2; \quad k \in \{\text{male, female}\} \end{aligned}$$

Numbers are fixed; randomization reveals one of the numbers.

Assume non-interference: each student's response depends only on that student's treatment.

If gender doesn't matter,

$$r_{ij\text{male}} = r_{ij\text{female}}.$$

Randomization

Conceptually, there are two levels of randomization:

- ① N_m students are randomly assigned to the male instructor, and the remaining N_f get the female instructor.
- ② Of the N_j assigned to instructor j , N_{jm} are told that the instructor is male, for $j = 1, 2$.

All $\binom{N_m}{N_{mm}} \times \binom{N_f}{N_{fm}}$ assignments of students to sections are equally likely.

Stratified two-sample test

- For each instructor, permute perceived gender assignments
- Use difference in mean ratings for female-identified minus male-identified

Stratified two-sample test

Results

In all categories, the male-identified instructor was rated higher.

Characteristic	M-F	perm P	t-test P
Overall	0.47	0.12	0.128
Caring	0.52	0.10	0.071
Consistent	0.47	0.21	0.045
Enthusiastic	0.57	0.06	0.112
Fair	0.76	0.01	0.188
Feedback	0.47	0.16	0.054
Helpful	0.46	0.17	0.049
Knowledgeable	0.35	0.29	0.038
Praise	0.67	0.01	0.153
Professional	0.61	0.07	0.124
Prompt	0.80	0.01	0.191
Respectful	0.61	0.06	0.124
Responsive	0.22	0.48	0.013

Omnibus Test

Nonparametric combination of tests (NPC): combine individual p-values into a single omnibus test when there are many responses

Test whether **all null** hypotheses are true or **at least one alternative** is true

Fisher's combining function

Let $\{P_j\}_{j=1}^J$ be p-values for J hypotheses. Define

$$X^2 = -2 \sum_{j=1}^J \ln(P_j)$$

If $\{P_j\}_{j=1}^J$ are independent and all nulls are true, then $X^2 \sim \chi_{2J}^2$.

Omnibus Test

Ratings by the same student for different categories are **dependent**.

⇒ Treat all ratings from a student as a vector and calibrate the distribution of X^2 using the this permutation distribution.

TO DO: CHECK THAT THIS IS LEGIBLE

NPC Permutation Procedure

- ① Calculate the vector of test statistics (use the **same permutation** of section memberships to compute all statistics), repeat a large number B times
- ② Compute the p-value for each individual variable in each permutation relative to the other values in the distribution
- ③ Apply the combining function to each vector of p-values.

Omnibus Test

```
# Initialize placeholders
ind = 0
test_distr = np.zeros( (10**5, len(categories)) )
pvalues = np.zeros( len(categories) )

# Loop over rating categories
for col in categories:
    (p, t, distr) = stratified_two_sample(
                    group=ratings.tagender,
                    response=ratings[col],
                    condition=ratings.taidgender,
                    alternative="two-sided",
                    stat="mean", seed = seed,
                    reps = 10**5, keep_dist = True)
    ind += 1
    test_distr[:,ind] = distr; pvalues[ind] = p

# NPC
omnibus_pvalue = npc(pvalues, test_distr, combine="fisher",
                      alternatives="two-sided")
```

Conclusions

Omnibus test: $P = 0$

- Reject the null hypothesis that there is no difference in ratings for any category
- The male-identified instructor was rated significantly higher than the female-identified instructor on several dimensions, **even on objective measures** such as how promptly assignments were returned
- SET measure something other than teaching effectiveness

NSGK Data

- Naomi Stark and Gilbert Kliman (NSGK) collected videos of therapy sessions with children on the autism spectrum
- A team of trained raters watched and tagged each 30-second interval of video from a collection of 183 clinically relevant tags
- Is tagging of therapist-patient interactions reliable (Millman et al. [2016])? Which tags do raters agree on?

Inter-rater reliability test

There are four dimensions. Can we simplify?

- Consider each clinical tag individually
- Do a partial hypothesis test for each video, then combine using NPC

NSGK	IRR
183 types of activity	T tags
8 videos	S strata
~ 35 segments/video	N_s items/stratum
10 raters	R raters

Inter-rater reliability test

Is agreement within columns better than expected by chance?

		Video segment			
		1	2	...	N_s
Rater	1	White	Dark	White	Dark
	2	Dark	White	White	Dark
	3	Dark	White	White	White
	4	White	White	White	White
	5	Dark	Dark	White	White
	6	White	Dark	White	White
	7	Dark	White	White	White
	8	White	Dark	White	White
	R	White	Dark	White	Dark

Inter-rater reliability test

Define

- $\{L_{s,i,r}\}$ = indicator for whether rater r tagged item i in stratum s
- $y_{si} = \sum_{r=1}^R L_{s,i,r}$ = number of raters who tagged item i in stratum s

The test statistic within stratum s is

$$\begin{aligned}\rho_s &\equiv \frac{1}{N_s \binom{R}{2}} \sum_{i=1}^{N_s} \sum_{r=1}^{R-1} \sum_{v=r+1}^R \mathbf{1}(L_{s,i,r} = L_{s,i,v}) \\ &= \frac{1}{N_s R(R-1)} \sum_{i=1}^{N_s} (y_{si}(y_{si} - 1) + (R - y_{si})(R - y_{si} - 1)).\end{aligned}$$

Inter-rater reliability test

Now we have a measure of concordance. What is the chance model?

		Video segment			
		1	2	...	N_s
Rater	1	White	Dark	White	Dark
	2	Dark	Dark	White	Dark
.		White	White	White	White
.		Dark	Dark	White	White
.		White	Dark	White	Dark
R		White	Dark	White	Dark

Permutation test

If tags are assigned completely at random, then

- any of 2^{N_s} assignments of tags are equally likely for each rater.
- raters assign tags independently of each other
- Each rater may have different “propensity” to assign a tag
 \Rightarrow condition on the number of items that a rater tagged.

IRR Permutation Test

- Compute ρ_s for $s = 1, \dots, S$ for the observed data
- Permute tags within rows, independently across rows and across strata, to get the permutation distributions of ρ_s , $s = 1, \dots, S$.
- For the observed and permutation ρ_s values, find the corresponding p-value P_s for each permutation, for $s = 1, \dots, S$.
- The NPC test statistic is $T = - \sum_{s=1}^S \frac{P_s}{\sqrt{N_s}}$.

Code

```
from permute.data import nsgk
from permute.irr import simulate_ts_dist, simulate_npc_dist

# load data, set video sizes
x = nsgk()
time_stamps = np.array([36, 32, 35, 37, 31, 35, 40, 32])

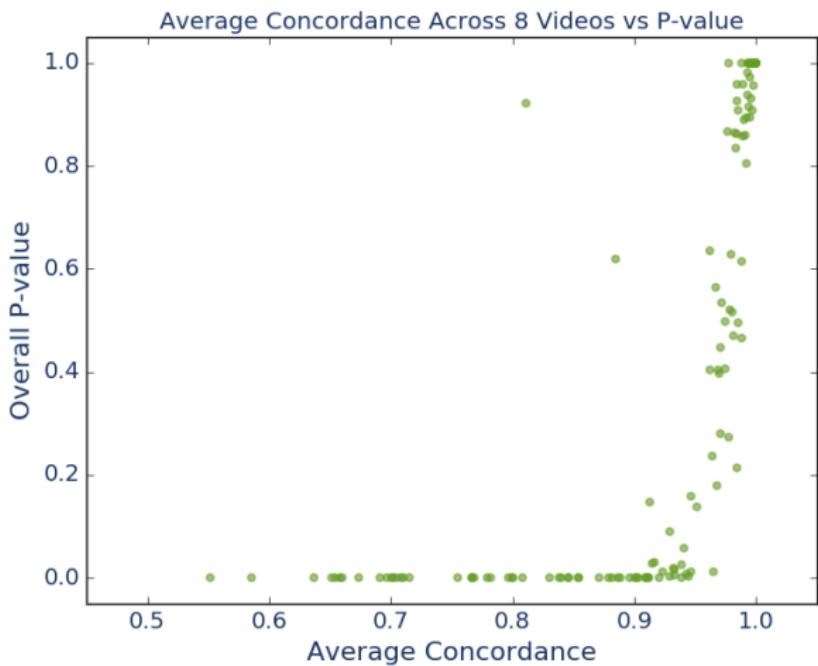
# Empty lists to store distrs and statistics for each video
d = []; tst = [] ; vid_temp = []

# Run analysis for a single category i
for j in range(len(x[i])): # loop over videos
    res = simulate_ts_dist(x[i][j], keep_dist=True)
    d.append(res["dist"])
    tst.append(res["obs_ts"])
    vid_temp.append(res["pvalue"])

# Combine permutation distributions for each video
perm_distr = np.asarray(d).transpose()
simulate_npc_dist(perm_distr, size=time_stamps,
                  obs_ts=tst, keep_dist=False)
```

Results

- 60 tags had $P < 0.05$
- Statistical vs practical significance – consult domain scientists
- Is there a more useful summary statistic than ρ_s ?



Reproducibility

Why should Statisticians worry about writing software?

- Ethics
- Impact

Monkey Cage

Does social science have a replication crisis?

RESEARCH ARTICLE

Estimating the reproducibility of psychological science

Open Science Collaboration^{*†}



OPEN ACCESS

ESSAY

Why Most Published Research Findings Are False

John P. A. Ioannidis

Published: August 30, 2005 • <http://dx.doi.org/10.1371/journal.pmed.0020124>

SundayReview

There's a replication crisis in biomedicine—and no one even knows how deep it runs.

By Daniel Engber

Why Do So Many Studies Fail to Replicate?

Gray Matter

By JAY VAN BAERL MAY 27, 2016

NATURE | EDITORIAL

Reality check on reproducibility

POLICY & ETHICS

Is There a Reproducibility Crisis in Science?

About 40% of economics experiments fail replication survey

By John Bohannon | Mar. 3, 2016, 2:00 PM

NATURE | NEWS

Over half of psychology studies fail reproducibility test

Largest replication study to date casts doubt on many published positive results.

Monya Baker

27 August 2015

Ethics

Much of the reproducibility crisis can be traced back to bad statistics.

- Publication bias: positive findings are more likely to get published
- P-hacking and the garden of forking paths (Gelman and Loken [2013])
- Inappropriate statistical tests (Randomization inference may help here)

We must make it easy for researchers to do the right statistics.

Impact

Let us own data science (Yu [2014]).

Statisticians have to build the tools to

- facilitate reproducible scientific research,
- enable people to use the methods we develop (correctly!), and
- influence the way people do statistics more broadly.

Download `permute`!

Permutation tests and confidence sets

[build](#) passing [coverage](#) 99%

Permutation tests and confidence sets for a variety of nonparametric testing and estimation problems, for a variety of randomization designs.

- **Website (including documentation):** <http://statlab.github.io/permute>
- **Mailing list:** <http://groups.google.com/group/permute>
- **Source:** <https://github.com/statlab/permute>
- **Bug reports:** <https://github.com/statlab/permute/issues>

Installation from binaries

```
$ pip install permute
```

<https://github.com/statlab/permute>

References

- Anne Boring, Kellie Ottoboni, and Philip B. Stark. Teaching evaluations (mostly) do not measure teaching effectiveness. *ScienceOpen Research*, 2016. doi: 10.14293/S2199-1006.1.SOR-EDU.AETBZC.v1.
- James V. Bradley. *Distribution-free statistical tests*. Prentice-Hall, 1968.
- Ronald A. Fisher. *Design of Experiments*. New York: Hafner, 1935.
- Andrew Gelman and Eric Loken. The garden of forking paths: Why multiple comparisons can be a problem, even when there is no "fishing expedition" or "p-hacking" and the research hypothesis was posited ahead of time. Unpublished paper, 2013.
- L. MacNell, A. Driscoll, and A. N. Hunt. What's in a name: Exposing gender bias in student ratings of teaching. *Innovative Higher Education*, pages 1–13, 2014.
- K. J. Millman, P. B. Stark, K. Ottoboni, and Naomi A. P. Stark. A case study in reproducible applied statistics: Is tagging of therapist-patient interactions reliable? Technical report, University of California, Berkeley, 2016. URL <https://github.com/statlab/nsgrk>.
- Bin Yu. Let us own data science. Institute of Mathematical Statistics (IMS) Presidential Address, ASC-IMS Joint Conference, Sydney, July 2014. URL <https://www.stat.berkeley.edu/~binyu/ps/papers2014/IMS-pres-address14-yu.pdf>.